

EDA supermarket case study

June 5, 2023

1 Description:

In this project, Python was used to do exploratory data analysis on a dataset from a supermarket. Significant conclusions were derived through investigating the data, cleaning the data, and doing statistical analysis and vizualisation. Actionable suggestions to enhance product ratings, pricing tactics, marketing initiatives and inventory control were developed as a result of these findings. The study shows how to find trends and make data-driven decisions for store improvement in a methodical manner.

1. Imported necessary libraries for data analysis and visualization in Python.
2. Imported the supermarket dataset and converted it into a data frame for further analysis.
3. Conducted an initial overview of the data to understand its structure and content.
4. Explored the dataset by examining the rows and columns, assessing its dimensions and characteristics.
5. Calculated statistical properties of each field, gaining insights into the distribution and variability of the data.
6. Determined the count of non-null values in each column and verified the data types for accurate analysis.
7. Plotted bar graphs to visualize columns with null values and stored the data frame into another variable for data manipulation and cleaning.
8. Replaced null values in the rating column with the mean of the existing ratings to ensure completeness and consistency.
9. Filled null values in the tax column by calculating 5% of the total amount, maintaining data integrity.
10. Replaced missing values in other columns (product line, gender, customer type, payment method) with the mode of their respective data to ensure representative values.
11. Imputed missing values in the branch column using city data, resulting in a complete and reliable dataset.
12. Modified the data types of the date and time columns to datetime format and added various features (e.g., month, year, days, weekend, quarter) using the date column.
13. Developed a function to analyze numeric columns (e.g., unit price, quantity) using swarm plots, strip plots, KDE plots, and volume plots to understand their spread and distribution.
14. Analyzed gender distribution through paragraph and pie charts, providing insights into customer demographics.
15. Created heatmaps to examine the relationships between numerical columns and identify correlations and patterns.
16. Generated pie charts to analyze sales across different stores and cities, allowing for market

performance assessment.

17. Implemented a comprehensive function to analyze total sales, mean sales, mean rating, and total quantity across categorical columns, providing insights into branch, city, customer type, time, gender, product line, and payment method.
18. Documented findings and actionable insights.

=====

1. Imported necessary libraries for data analysis and visualization in Python.

```
[78]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

2. Imported the supermarket dataset and converted it into a data frame for further analysis.

```
[79]: df=pd.read_csv('D:/DS/resume projects/supermarket EDA/columns_description.csv')
dfdict=pd.read_csv('D:/DS/resume projects/supermarket EDA/data_dict.csv')
```

3. Conducted an initial overview of the data to understand its structure and content.

```
[80]: df.head()
```

```
[80]:      Invoice ID Branch CustomerID      City Customer type Gender \
0  750-67-8428      A      C1888      Yangon      Member  Female
1  226-31-3081      C      C1475  Naypyitaw      Normal  Female
2  631-41-3108      A      C1746      Yangon      Normal   Male
3  123-19-1176      A      C1896      Yangon      Member   Male
4  373-73-7910      A      C1790      Yangon      Normal   Male

      Product line  Unit price  Quantity  Tax 5%  Total      Date \
0  Health and beauty      74.69         10  37.3450  746.90  21-02-2019
1  Health and beauty      15.28          6   4.5840   91.68  27-05-2019
2  Health and beauty      46.33          7  16.2155  324.31  27-12-2019
3  Health and beauty      58.22         11  32.0210  640.42  15-11-2019
4  Health and beauty      86.31          7  30.2085  604.17  31-03-2019

      Time      Payment      cogs  gross margin percentage  gross income \
0  13:08      Ewallet  711.333333          4.761905          35.566667
1  10:29        Cash    76.400000          4.761905          15.280000
2  13:23  Credit card  324.310000          4.761905           0.000000
3  20:33      Ewallet  465.760000          4.761905          174.660000
4  10:37      Ewallet  604.170000          4.761905           0.000000

      Rating  Longitude  Latitude
0         9.1     96.1735    16.8409
1        10.0     96.0785    19.7633
2         7.4     96.1735    16.8409
3         8.4     96.1735    16.8409
```

```
4      NaN      96.1735      16.8409
```

```
[81]: dfdict
```

```
[81]:
```

	Field	Description
0	Invoice ID	Invoice ID of the transaction
1	Branch	One out of 3 branches. Every city belongs to a...
2	CustomerID	Customer ID of the customer doing transaction
3	City	City where the tx took place. The chain has st...
4	Customer Type	Where a member or normal customer
5	Gender	Male or Female
6	Product Line	Product line of the product purchased
7	Unit Price	Unit price of product purchased
8	Quantity	Qty purchased
9	Tax 5%	Tax as a fixed % of invoice
10	Total	Total Invoice amount
11	Date	Date of tx
12	Time	Time of tx
13	Payment	Payment mode
14	cogs	Cost of goods sold
15	gross margin	Margin %
16	gross income	Margin amount
17	Rating	Rating given by cx out of 10
18	Longitude	Geo field for tx location
19	Latitude	Geo field for tx location

```
[116]: df.style.background_gradient(cmap='cool')
```

```
<pandas.io.formats.style.Styler object at 0x000001FF980652D0>
```

4. Explored the dataset by examining the rows and columns, assessing its dimensions and characteristics.

```
[83]: l=df.shape
print('The supermak=rket data has',l[0], 'rows and',l[1], 'columns')
```

```
The supermak=rket data has 1000 rows and 20 columns
```

5. Calculated statistical properties of each field, gaining insights into the distribution and variability of the data.

```
[84]: df.describe().style.background_gradient(cmap='hot')
```

```
[84]: <pandas.io.formats.style.Styler at 0x1ff97d7d030>
```

6. Determined the count of non-null values in each column and verified the data types for accurate analysis.

```
[85]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Invoice ID                            1000 non-null   object
1   Branch                               806 non-null    object
2   CustomerID                           1000 non-null   object
3   City                                  1000 non-null   object
4   Customer type                         952 non-null    object
5   Gender                               975 non-null    object
6   Product line                         977 non-null    object
7   Unit price                           1000 non-null   float64
8   Quantity                             1000 non-null   int64
9   Tax 5%                               896 non-null    float64
10  Total                                1000 non-null   float64
11  Date                                 1000 non-null   object
12  Time                                 1000 non-null   object
13  Payment                             979 non-null    object
14  cogs                                1000 non-null   float64
15  gross margin percentage              1000 non-null   float64
16  gross income                        1000 non-null   float64
17  Rating                              857 non-null    float64
18  Longitude                            1000 non-null   float64
19  Latitude                             1000 non-null   float64
dtypes: float64(9), int64(1), object(10)
memory usage: 156.4+ KB

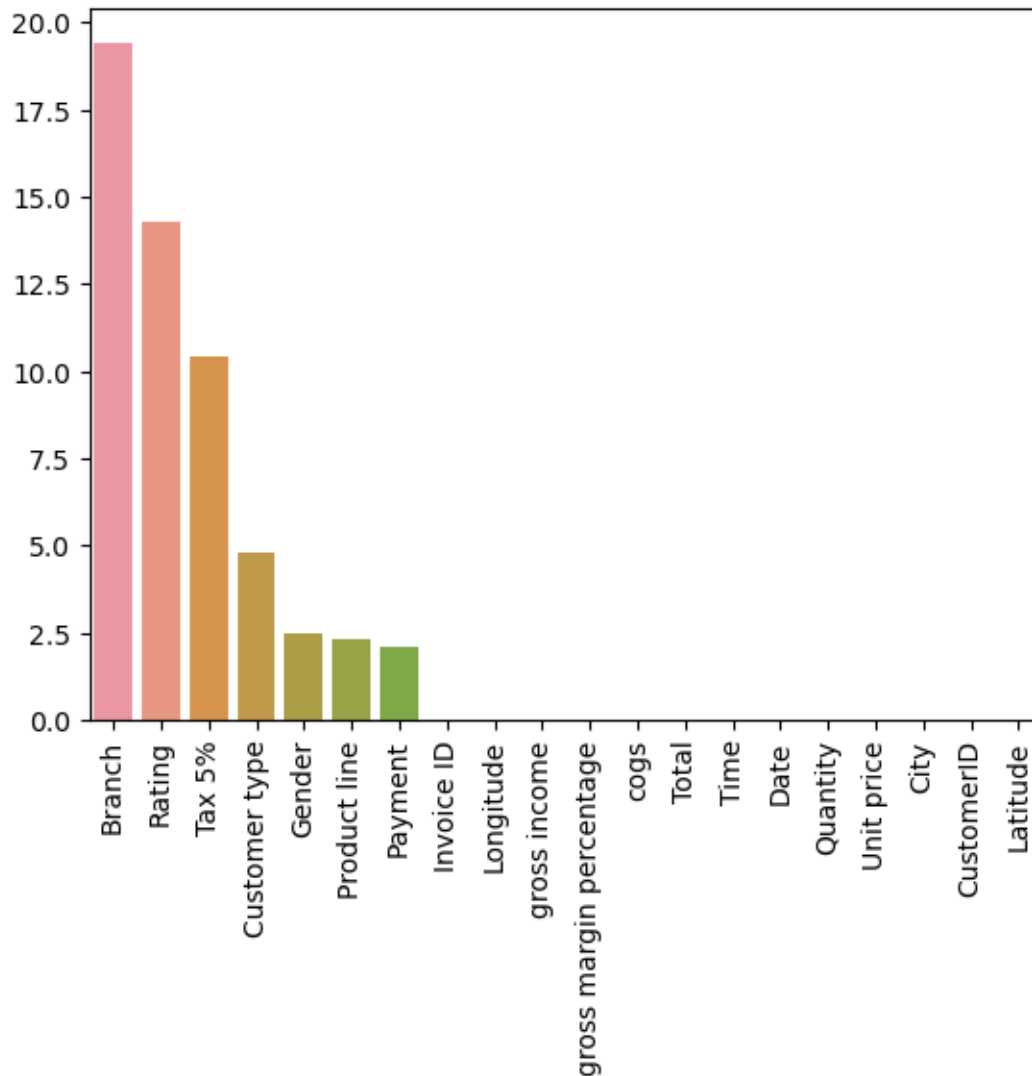
```

7. Plotted bar graphs to visualize columns with null values and stored the data frame into another variable for data manipulation and cleaning.

```

[86]: f=(df.isna().mean()*100).sort_values(ascending=False)
      sns.barplot(x=f.index,y=f.values)
      plt.xticks(rotation=90)
      plt.show()

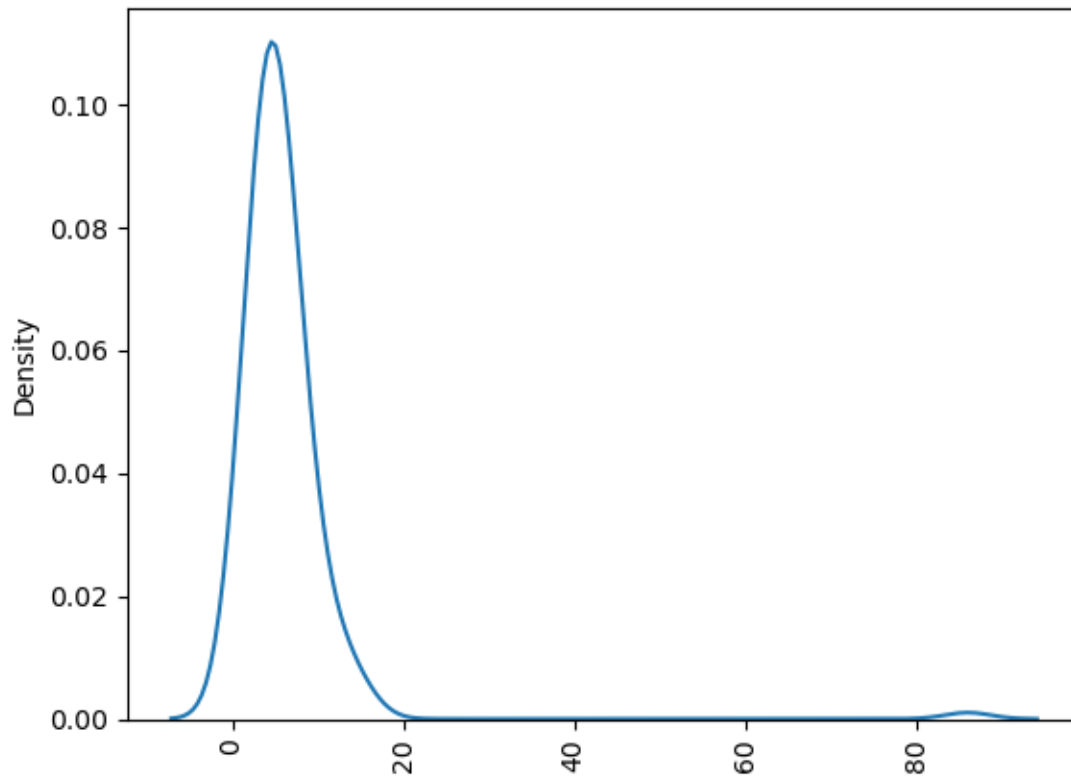
```



```
[87]: #original data in df and manipulation in sdf
sdf=df
```

8. Replaced null values in the rating column with the mean of the existing ratings to ensure completeness and consistency.

```
[88]: #Rating is summeticaly distributed replacing null with mean
b=df.Rating.value_counts()
sns.kdeplot(b.values)
plt.xticks(rotation=90)
plt.show()
```



```
[89]: df.Rating.describe()
```

```
[89]: count      857.000000  
      mean        7.462625  
      std        1.776179  
      min         4.000000  
      25%        5.900000  
      50%        7.455000  
      75%        9.100000  
      max       10.000000  
      Name: Rating, dtype: float64
```

```
[90]: sdf['Rating']=df['Rating'].fillna(df.Rating.mean())
```

```
[91]: sdf.Rating.describe()
```

```
[91]: count      1000.000000  
      mean        7.462625  
      std        1.644148  
      min         4.000000  
      25%        6.200000  
      50%        7.462625
```

```
75%          8.820000
max          10.000000
Name: Rating, dtype: float64
```

9. Filled null values in the tax column by calculating 5% of the total amount, maintaining data integrity.

```
[92]: # filling null of tax column by calculating it from total column
sdf['Tax 5%']=df['Tax 5%'].fillna(df['Total']*.05)
```

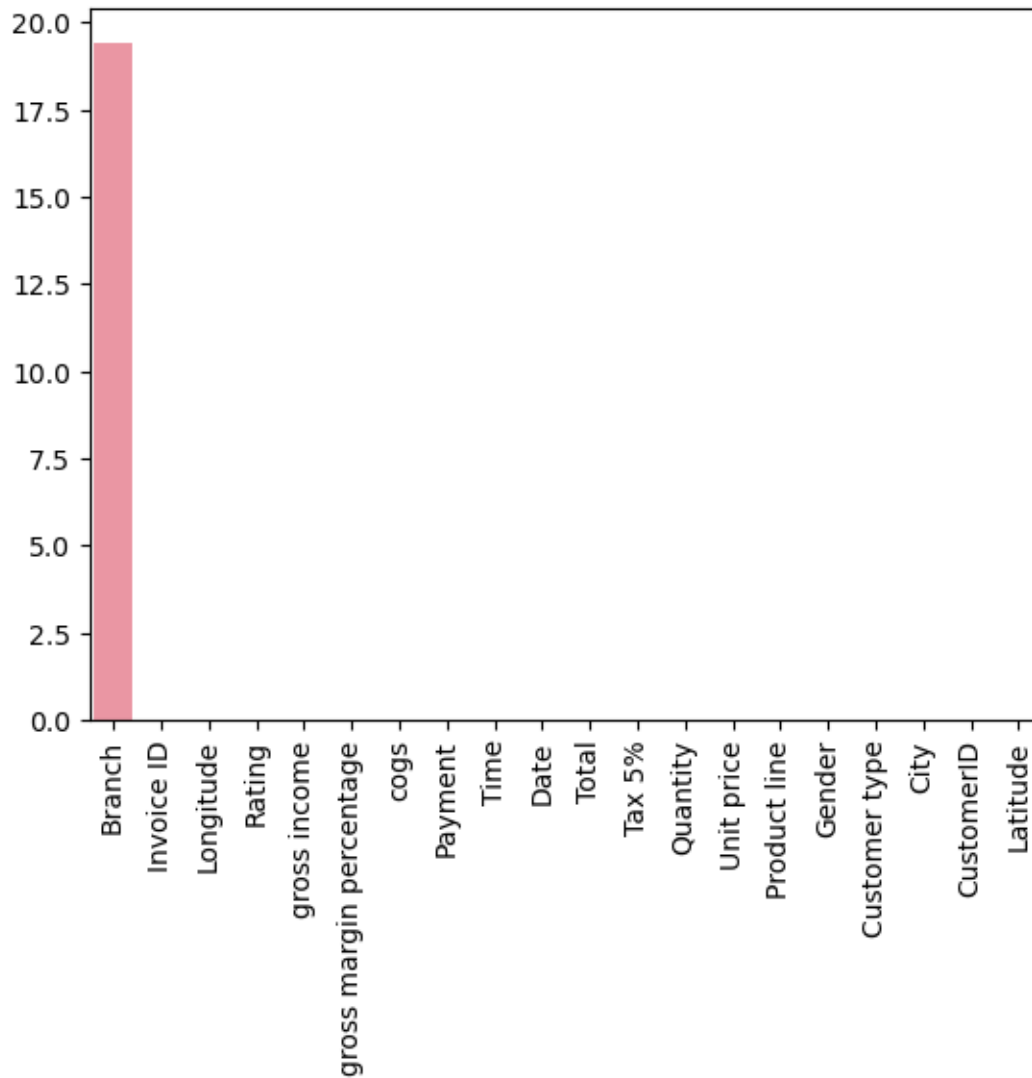
10. Replaced missing values in other columns (product line, gender, customer type, payment method) with the mode of their respective data to ensure representative values.

```
[93]: # filling null values in categorical column by mode
def repmode(col):
    sdf[col]=df[col].fillna(df[col].mode()[0])

repmode('Product line')
repmode('Gender')
repmode('Customer type')
repmode('Payment')
```

11. Imputed missing values in the branch column using city data, resulting in a complete and reliable dataset.

```
[94]: # Branch column need to be cleaned
f=(df.isna().mean()*100).sort_values(ascending=False)
sns.barplot(x=f.index,y=f.values)
plt.xticks(rotation=90)
plt.show()
```



```
[95]: sdf.groupby(['City', 'Branch'])['Invoice ID'].count()
```

```
[95]: City      Branch
Mandalay    B         265
Naypyitaw   C         260
Yangon      A         281
Name: Invoice ID, dtype: int64
```

```
[96]: sdf.loc[(sdf['Branch'].isnull() & (sdf['City'] == 'Mandalay'), 'Branch'] = 'B'
sdf.loc[(sdf['Branch'].isnull() & (sdf['City'] == 'Naypyitaw'), 'Branch'] = 'C'
sdf.loc[(sdf['Branch'].isnull() & (sdf['City'] == 'Yangon'), 'Branch'] = 'A'
```

```
[97]: sdf.isna().mean()*100
```



```
[97]: Invoice ID          0.0
      Branch            0.0
      CustomerID        0.0
      City              0.0
      Customer type      0.0
      Gender            0.0
      Product line       0.0
      Unit price         0.0
      Quantity          0.0
      Tax 5%            0.0
      Total             0.0
      Date              0.0
      Time              0.0
      Payment           0.0
      cogs              0.0
      gross margin percentage 0.0
      gross income       0.0
      Rating            0.0
      Longitude         0.0
      Latitude          0.0
      dtype: float64
```

12. Modified the data types of the date and time columns to datetime format and added various features (e.g., month, year, days, weekend, quarter) using the date column.

```
[98]: sdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Invoice ID            1000 non-null  object
1   Branch               1000 non-null  object
2   CustomerID           1000 non-null  object
3   City                 1000 non-null  object
4   Customer type        1000 non-null  object
5   Gender               1000 non-null  object
6   Product line         1000 non-null  object
7   Unit price           1000 non-null  float64
8   Quantity             1000 non-null  int64
9   Tax 5%               1000 non-null  float64
10  Total                1000 non-null  float64
11  Date                 1000 non-null  object
12  Time                 1000 non-null  object
13  Payment              1000 non-null  object
14  cogs                 1000 non-null  float64
15  gross margin percentage 1000 non-null  float64
```

```

16 gross income          1000 non-null    float64
17 Rating                1000 non-null    float64
18 Longitude             1000 non-null    float64
19 Latitude              1000 non-null    float64
dtypes: float64(9), int64(1), object(10)
memory usage: 156.4+ KB

```

```

[99]: #correcting datatype of time,date column
sdf['Date']=sdf['Date'].astype('datetime64')
sdf['Time']=sdf['Time'].astype('datetime64')
sdf.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Invoice ID             1000 non-null  object
1   Branch                1000 non-null  object
2   CustomerID            1000 non-null  object
3   City                  1000 non-null  object
4   Customer type         1000 non-null  object
5   Gender                1000 non-null  object
6   Product line          1000 non-null  object
7   Unit price            1000 non-null  float64
8   Quantity              1000 non-null  int64
9   Tax 5%                1000 non-null  float64
10  Total                 1000 non-null  float64
11  Date                  1000 non-null  datetime64[ns]
12  Time                  1000 non-null  datetime64[ns]
13  Payment               1000 non-null  object
14  cogs                  1000 non-null  float64
15  gross margin percentage 1000 non-null  float64
16  gross income          1000 non-null  float64
17  Rating                1000 non-null  float64
18  Longitude             1000 non-null  float64
19  Latitude              1000 non-null  float64
dtypes: datetime64[ns](2), float64(9), int64(1), object(8)
memory usage: 156.4+ KB

```

```

C:\Users\Kundan Mourya\AppData\Local\Temp\ipykernel_19992\846853813.py:2:
UserWarning: Parsing dates in DD/MM/YYYY format when dayfirst=False (the
default) was specified. This may lead to inconsistently parsed dates! Specify a
format to ensure consistent parsing.
sdf['Date']=sdf['Date'].astype('datetime64')

```

```

[100]: sdf['Month']=sdf['Date'].dt.month
sdf['Year']=sdf['Date'].dt.year

```

```

sdf['Day']=sdf['Date'].dt.day
sdf['Weekday']=sdf['Date'].dt.day_name()
sdf['Quarter']=sdf['Date'].dt.quarter
#df['week_number'] = df['Date'].dt.week
df['annual_week_number'] = df['Date'].dt.isocalendar().week

sdf['Hour']=sdf['Time'].dt.hour
sdf['Minute']=sdf['Time'].dt.minute

```

[101]: sdf.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 28 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Invoice ID                            1000 non-null   object
1   Branch                               1000 non-null   object
2   CustomerID                           1000 non-null   object
3   City                                  1000 non-null   object
4   Customer type                         1000 non-null   object
5   Gender                                1000 non-null   object
6   Product line                          1000 non-null   object
7   Unit price                            1000 non-null   float64
8   Quantity                              1000 non-null   int64
9   Tax 5%                                1000 non-null   float64
10  Total                                 1000 non-null   float64
11  Date                                  1000 non-null   datetime64[ns]
12  Time                                  1000 non-null   datetime64[ns]
13  Payment                              1000 non-null   object
14  cogs                                  1000 non-null   float64
15  gross margin percentage               1000 non-null   float64
16  gross income                          1000 non-null   float64
17  Rating                                1000 non-null   float64
18  Longitude                             1000 non-null   float64
19  Latitude                              1000 non-null   float64
20  Month                                 1000 non-null   int64
21  Year                                  1000 non-null   int64
22  Day                                   1000 non-null   int64
23  Weekday                               1000 non-null   object
24  Qruarter                              1000 non-null   int64
25  annual_week_number                    1000 non-null   UInt32
26  Hour                                  1000 non-null   int64
27  Minute                                1000 non-null   int64
dtypes: UInt32(1), datetime64[ns](2), float64(9), int64(7), object(9)
memory usage: 215.9+ KB

```

[102]: sdf.head()

```
[102]: Invoice ID Branch CustomerID City Customer type Gender \
0 750-67-8428 A C1888 Yangon Member Female
1 226-31-3081 C C1475 Naypyitaw Normal Female
2 631-41-3108 A C1746 Yangon Normal Male
3 123-19-1176 A C1896 Yangon Member Male
4 373-73-7910 A C1790 Yangon Normal Male

Product line Unit price Quantity Tax 5% ... Longitude Latitude \
0 Health and beauty 74.69 10 37.3450 ... 96.1735 16.8409
1 Health and beauty 15.28 6 4.5840 ... 96.0785 19.7633
2 Health and beauty 46.33 7 16.2155 ... 96.1735 16.8409
3 Health and beauty 58.22 11 32.0210 ... 96.1735 16.8409
4 Health and beauty 86.31 7 30.2085 ... 96.1735 16.8409

Month Year Day Weekday Qruarter annual_week_number Hour Minute
0 2 2019 21 Thursday 1 8 13 8
1 5 2019 27 Monday 2 22 10 29
2 12 2019 27 Friday 4 52 13 23
3 11 2019 15 Friday 4 46 20 33
4 3 2019 31 Sunday 1 13 10 37
```

[5 rows x 28 columns]

```
[103]: sdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 28 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Invoice ID                            1000 non-null   object
1   Branch                               1000 non-null   object
2   CustomerID                           1000 non-null   object
3   City                                  1000 non-null   object
4   Customer type                         1000 non-null   object
5   Gender                               1000 non-null   object
6   Product line                         1000 non-null   object
7   Unit price                           1000 non-null   float64
8   Quantity                             1000 non-null   int64
9   Tax 5%                              1000 non-null   float64
10  Total                                1000 non-null   float64
11  Date                                 1000 non-null   datetime64[ns]
12  Time                                 1000 non-null   datetime64[ns]
13  Payment                              1000 non-null   object
14  cogs                                 1000 non-null   float64
15  gross margin percentage              1000 non-null   float64
16  gross income                         1000 non-null   float64
17  Rating                              1000 non-null   float64
```

18	Longitude	1000 non-null	float64
19	Latitude	1000 non-null	float64
20	Month	1000 non-null	int64
21	Year	1000 non-null	int64
22	Day	1000 non-null	int64
23	Weekday	1000 non-null	object
24	Quarter	1000 non-null	int64
25	annual_week_number	1000 non-null	UInt32
26	Hour	1000 non-null	int64
27	Minute	1000 non-null	int64

dtypes: UInt32(1), datetime64[ns](2), float64(9), int64(7), object(9)
memory usage: 215.9+ KB

13. Developed a function to analyze numeric columns (e.g., unit price, quantity) using swarm plots, strip plots, KDE plots, and volume plots to understand their spread and distribution.

```
[104]: def fun(col):
        fig, ax = plt.subplots(4, 1, figsize=(7, 5))
        sns.swarmplot(data=sdf, x=col, ax=ax[0])
        ax[0].set_title('swarmplot of ' + col)

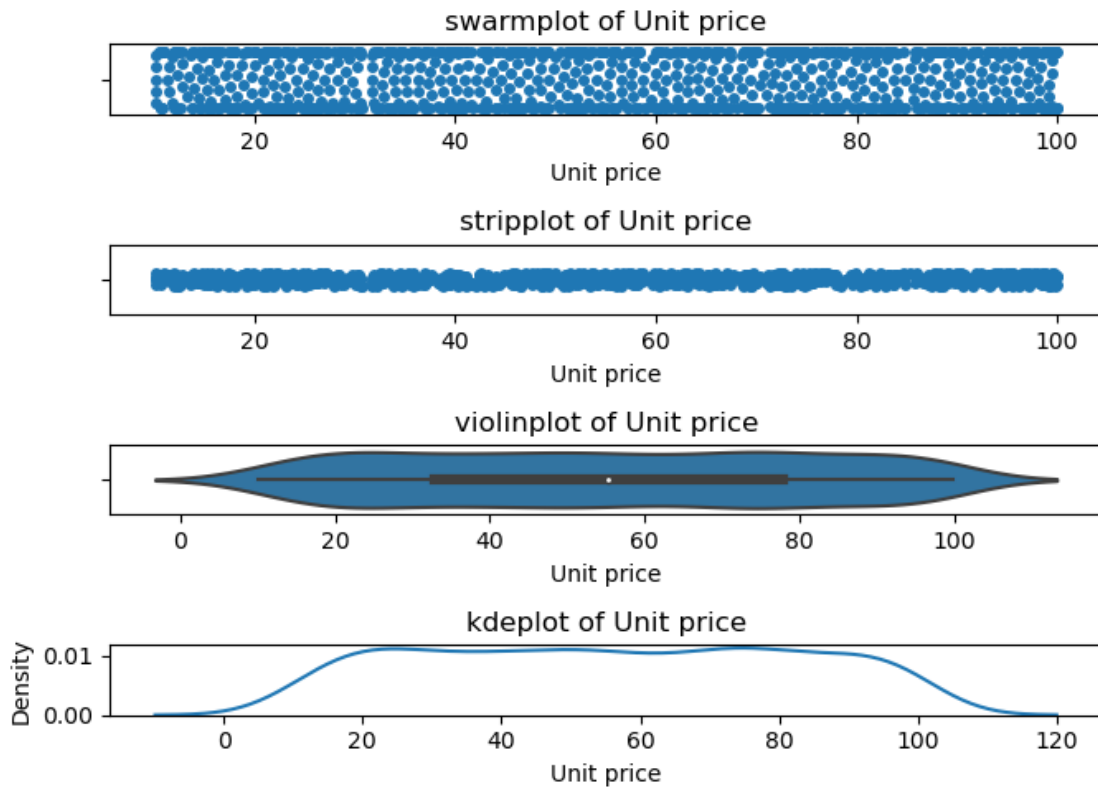
        sns.stripplot(data=sdf, x=col, ax=ax[1])
        ax[1].set_title('stripplot of ' + col)

        sns.violinplot(data=sdf, x=col, ax=ax[2])
        ax[2].set_title('violinplot of ' + col)

        sns.kdeplot(sdf[col], ax=ax[3])
        ax[3].set_title('kdeplot of ' + col)
        plt.tight_layout()
        plt.show()

        fun('Unit price')
```

C:\Users\Kundan Mourya\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 64.4% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
warnings.warn(msg, UserWarning)



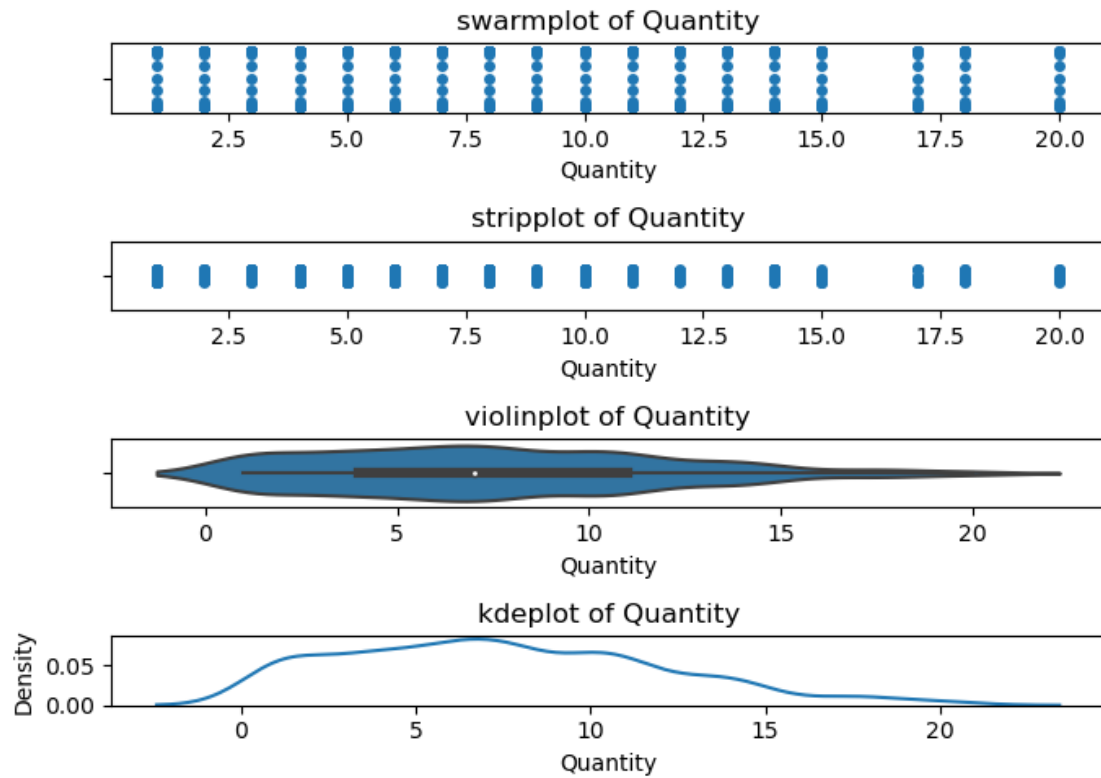
```
[105]: fun('Quantity')
```

C:\Users\Kundan Mourya\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 7.1% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.

```
warnings.warn(msg, UserWarning)
```

C:\Users\Kundan Mourya\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 91.0% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.

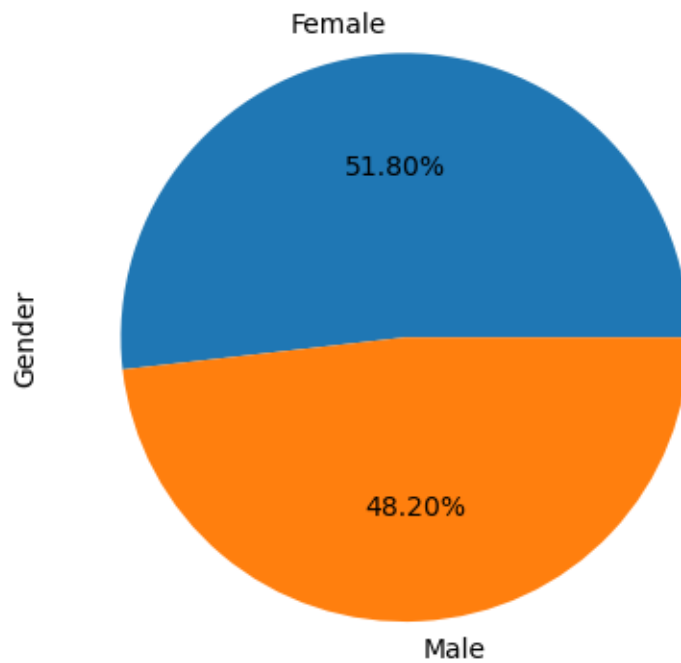
```
warnings.warn(msg, UserWarning)
```



14. Analyzed gender distribution through pie charts, providing insights into customer demographics.

```
[106]: sdf.Gender.value_counts().plot.pie(autopct='%1.2f%%')
```

```
[106]: <Axes: ylabel='Gender'>
```



15. Created heatmaps to examine the relationships between numerical columns and identify correlations and patterns.

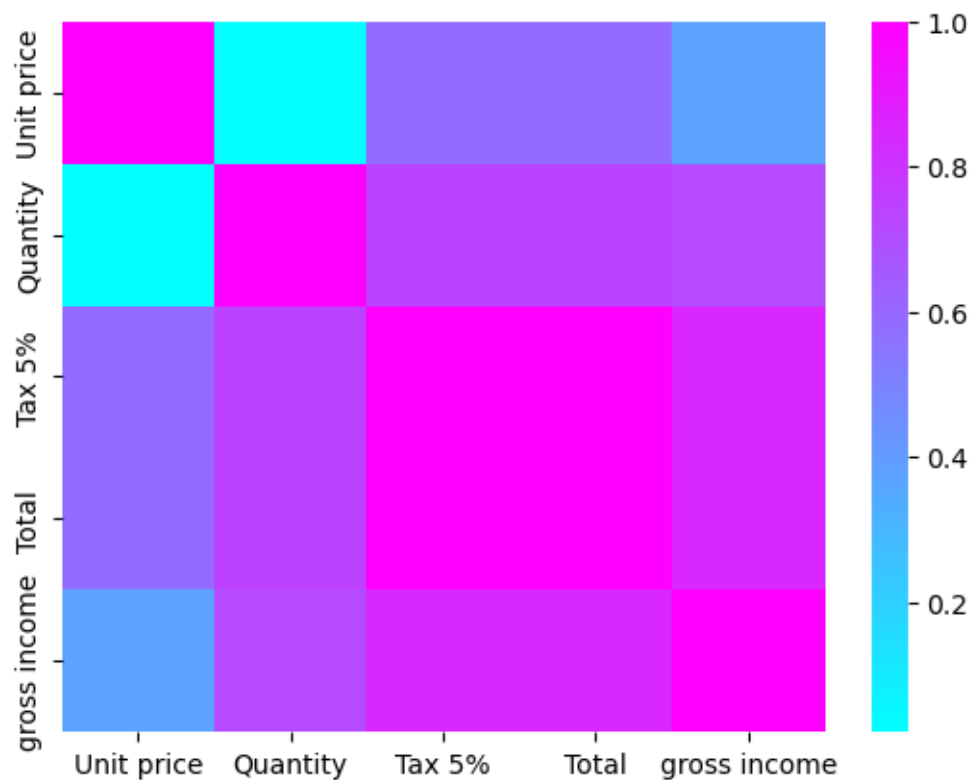
```
[107]: sdf[['City', 'Longitude', 'Latitude']].drop_duplicates()
```

```
[107]:
```

	City	Longitude	Latitude
0	Yangon	96.1735	16.8409
1	Naypyitaw	96.0785	19.7633
9	Mandalay	96.0891	21.9588

```
[108]: sns.heatmap(sdf[['Unit price', 'Quantity', 'Tax 5%', 'Total', 'gross income']].  
               ↪corr(), cmap='cool')
```

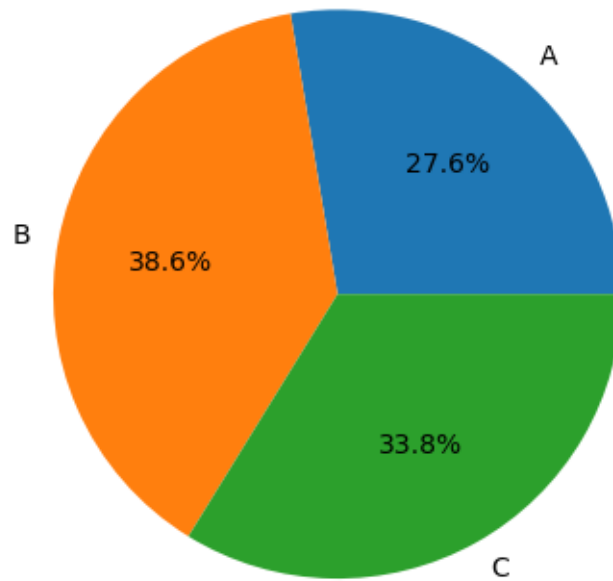
```
[108]: <Axes: >
```

16. Generated pie charts to analyze sales across different stores/cities, allowing for market performance assessment.

```
[109]: h=sdf.groupby(['Branch'])['Total'].sum()
plt.pie(h.values,labels=h.index,autopct='%1.1f%%')
plt.title('Toatal sales from each branch')
plt.show()
```

Toatal sales from each branch



17. Implemented a comprehensive function to analyze total sales, mean sales, mean rating, and total quantity across categorical columns, providing insights into branch, city, customer type, time, gender, product line, and payment method.

```
[110]: catcol=list(sdf.columns[sdf.dtypes=='O'])
catcol.remove('Invoice ID')
catcol.remove('CustomerID')
catcol
```

```
[110]: ['Branch',
        'City',
        'Customer type',
        'Gender',
        'Product line',
        'Payment',
        'Weekday']
```

```
[111]: def ubit(col):

        a=(sdf.groupby(sdf[col])['Total'].sum()).sort_values(ascending=False)
        b=(sdf.groupby([sdf[col]])['Total'].mean()).sort_values(ascending=False)
        c=(sdf.groupby([sdf[col]])['Rating'].mean()).sort_values(ascending=False)
        d=(sdf.groupby([sdf[col]])['Quantity'].sum()).sort_values(ascending=False)
```

```

fig,ax=plt.subplots(2,2,figsize=(15,5))

sns.lineplot(x=a.index, y=a.values, ax=ax[0, 0])
ax[0,0].set_title('Total sales by '+col)

ax[0,1].bar(height=b.values,x=b.index)
ax[0,1].set_title('Mean sales by '+col)

ax[1,0].barh(y=c.index, width=c.values)
ax[1,0].set_title('Mean Rating by '+col)

ax[1,1].barh(y=d.index, width=d.values)
ax[1,1].set_title('Total Quantity by '+col)

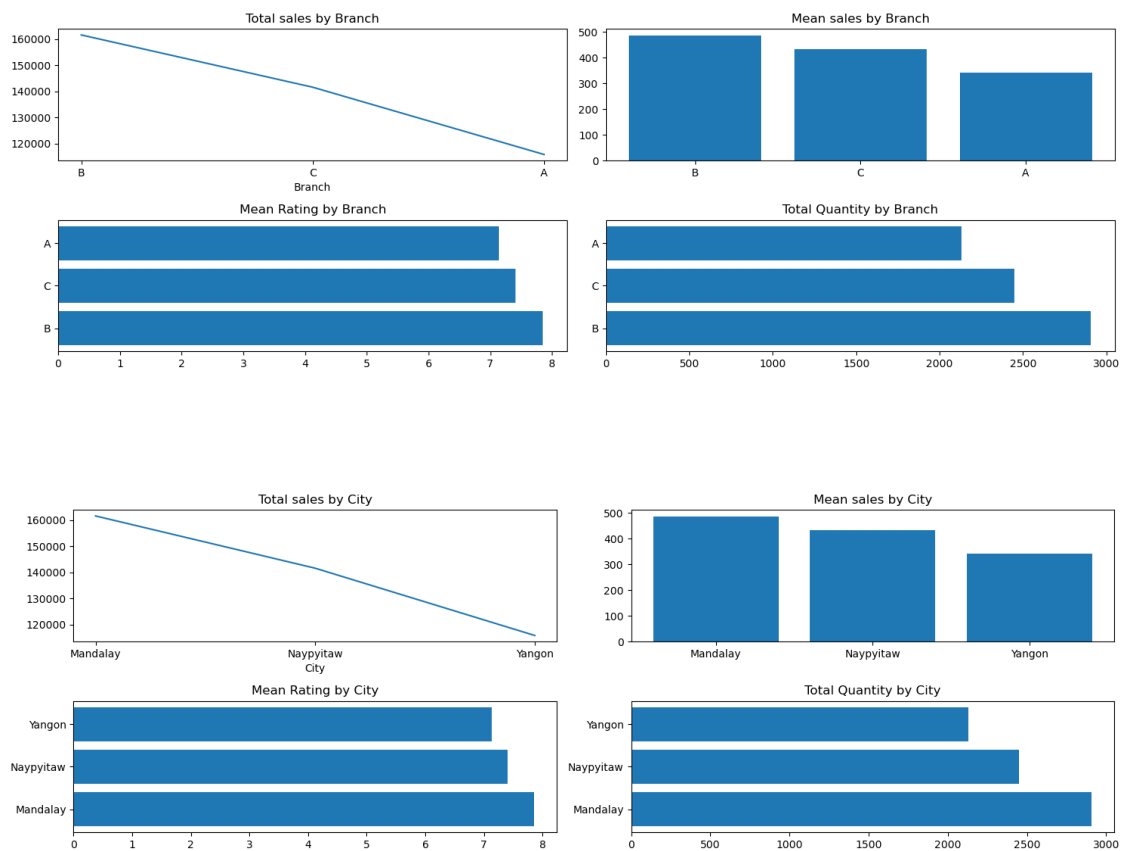
plt.tight_layout()

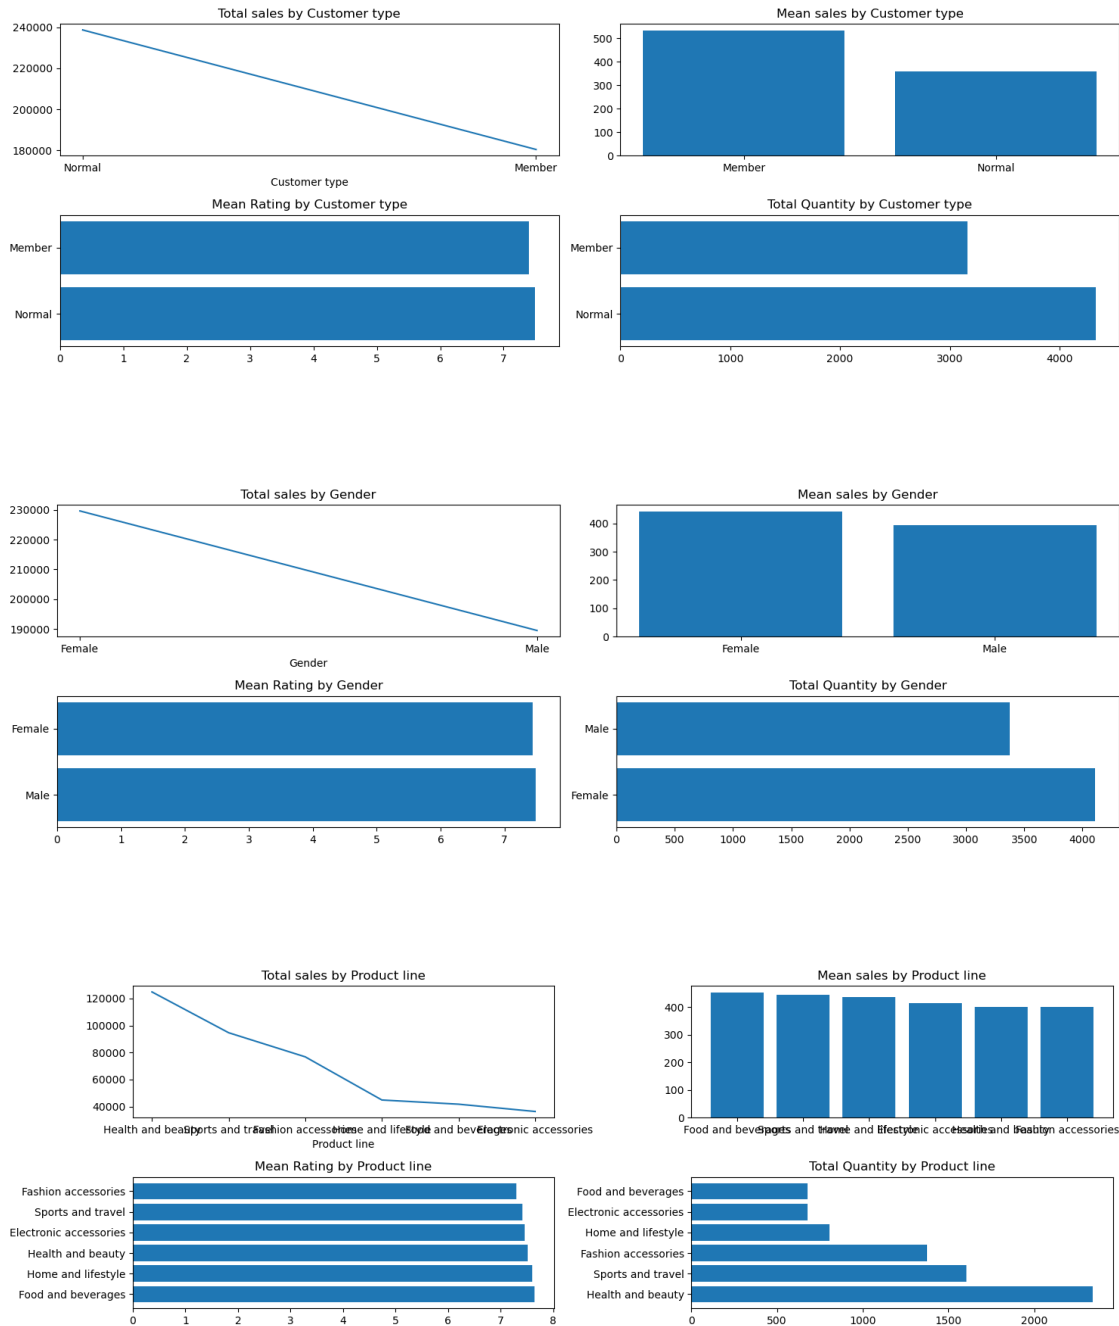
```

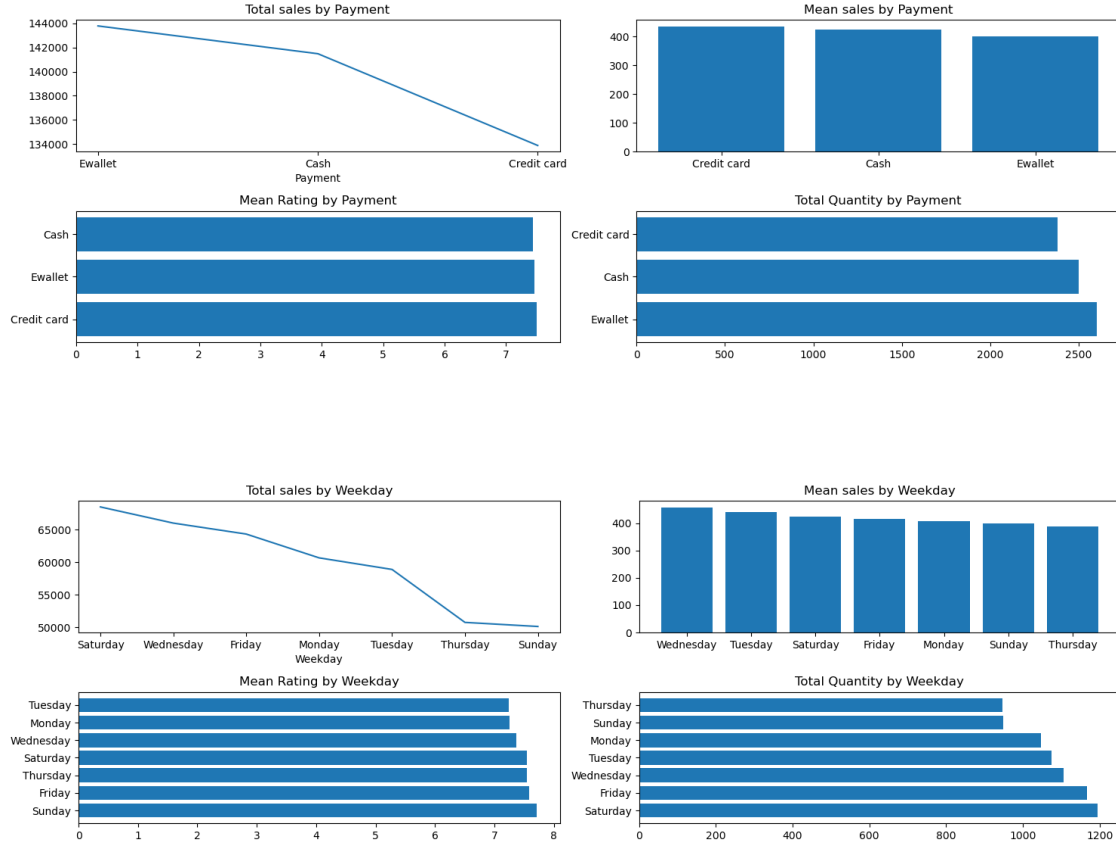
```

[112]: for i in catcol:
        ubit(i)

```







2 Insights

- **Improve product ratings:** With the knowledge that missing ratings were replaced with the mean rating, the supermarket can focus on improving the quality and customer happiness of items with lower ratings in order to better the overall shopping experience for customers.
- **Modify pricing strategies:** By evaluating the variation in unit prices and quantities, the supermarket may spot patterns in pricing and make necessary price adjustments. Reevaluating pricing tactics for particular products, for instance, may be advantageous if their prices vary widely.
- **Boost marketing efforts:** Bar graphs showing gender distribution can be utilised for more accurate marketing activities. In order to increase consumer engagement and sales, the supermarket can customise promotional events, product displays, and adverts to individual gender preferences.
- **The supermarket can pinpoint failing locations and take action to improve inventory levels, product selection, and operational efficiency in those regions through analysing sales distribution across different stores and cities.**
- **Optimising product line offerings:** Supermarkets may determine the best-performing product categories and make informed decisions about inventory management, marketing strategies,

and possible growth of profitable product lines via the insights gained from analysing the total sales, mean sales, mean rating, and total quantity across different product lines.

- Determine peak sales times: The supermarket can identify periods of highest sales and make suitable preparations by using the data gathered from the date column, such as month, year, days, weekend, and quarter. To maximise sales at times of high demand, this could involve altering employee numbers, inventory restocking, and promotional efforts.

2.0.1 Based on the data that was gathered, these choices might result in greater customer happiness, more sales, improved management of inventory, and overall company success for the supermarket.